

EU's AI Act

João Cortes & Michelle Wang

“AI should be a tool for people and be a force for good in society with the ultimate aim of increasing human well-being.”

— AI Act

Context

The European Commission has been the first major regulator to craft a law on Artificial Intelligence on the 21st of April 2021.

This Act follows from the release of a White Paper on AI in 2020 – developed after a public consultation on this matter.

After multiple rounds of revision between EU Presidencies and comments received by many affected stakeholders, the EU AI Act was published in the Official Journal (OJ) of the European Union on 12 July 2024 and entered into force on August 2024.

The final version of the document is a comprehensive regulation that spans 180 recitals and 113 Articles. (144 pages long)

Development Procedure



April 21, 2021

Initial Proposal Published

The European Commission releases the first regulatory framework proposal for AI.

Dec 2022

Council "General Approach"

The Council of the EU adopts its negotiating position after multiple compromise texts.

Dec 2023

Provisional Agreement Reached

Parliament and Council reach a landmark political deal on the AI Act's final rules.

May 2024

Formal Council Adoption

The European Council gives the final green light, concluding the legislative process.

July 12, 2024

Publication in Official Journal

The Act is published in the OJ of the EU, setting the stage for entry into force.

Logic

1. What is being regulated?
- 2. Why is it risky?
- 3. Who is responsible?
- 4. What obligations apply?
- 5. What remains unresolved?

Initial Objective of the Act

Defined Need

Addressing the opacity, complexity, bias, and partially autonomous behavior of certain AI systems to ensure:

- Compatibility with fundamental rights
- Facilitation of legal rules enforcement

Core Objectives

- **Safety & Rights:** Ensure AI systems respect fundamental rights and Union values.
- **Legal Certainty:** Facilitate investment and innovation through clear rules.
- **Market Development:** Create a single market for lawful, safe, and trustworthy AI.
- **Fragmentation:** Prevent market fragmentation across the Union.

A Risk Level-based Assessment

Risk-Based Obligations

While low-risk apps are unregulated, medium/high-risk systems require compulsory conformity self-assessment before going to market.

Prohibited Practices

The proposal focuses on the prohibition of biometric identification and social scoring.

Governance & Expertise

Integration of an European AI Board to support the Commission with expertise.


Support for Innovation

Authorities may establish AI regulatory sandboxes for development and testing of innovative systems.

The Risk Pyramid (The Core Framework)

A Risk-based Approach

The AI Act defines 4 levels of risk for AI systems:



A Risk-based Approach

The AI Act defines 4 levels of risk for AI systems:

UNACCEPTABLE RISK:

Strictly prohibited practices that threaten safety and rights.

HIGH RISK:

Systems subject to strict compliance and human oversight.

LIMITED RISK:

AI systems with specific transparency obligations. Includes chatbots and deep fakes.

MINIMAL RISK:

Unregulated applications like AI-enabled video games. Includes most applications.

The Risk Pyramid (The Core Framework)

Discussion:

Try to list 1-2 examples for each of the categories

1. Unacceptable risk
2. High risk
3. Limited risk
4. Minimal risk

A Risk-based Approach

The AI Act defines 4 levels of risk for AI systems:

UNACCEPTABLE RISK:

Strictly prohibited practices that threaten safety and rights.

HIGH RISK:

Systems subject to strict compliance and human oversight.

LIMITED RISK:

AI systems with specific transparency obligations. Includes chatbots and deep fakes.

MINIMAL RISK:

Unregulated applications like AI-enabled video games. Includes most applications.

1. Unacceptable AI Risk Practices



Cognitive Manipulation

AI systems that use subliminal techniques to distort behavior and cause physical or psychological harm.



Social Scoring

Classifying individuals based on social behavior or personality traits leading to unfavorable treatment.



Real-time Biometrics

Remote biometric identification in publicly accessible spaces for law enforcement (with limited exceptions).



Exploiting Vulnerabilities

Targeting specific groups based on age, disability, or socio-economic situation to distort behavior.

CONSEQUENCES

These practices are strictly forbidden under the EU AI Act due to the threat they pose to safety and fundamental rights.



Note: The EU AI Act does not ban biometrics entirely. It bans or restricts specific biometric uses when they threaten fundamental rights, especially in public space, policing, workplaces, education, and sensitive categorisation.

1. Unacceptable AI Risk Practices

Article 5: Prohibited AI Practices

Article 5(1)(a): harmful **manipulation**, deception, or subliminal techniques.

Article 5(1)(b): **exploitation of vulnerabilities** due to age, disability, or social/economic situation.

Article 5(1)(c): **social scoring**.

Article 5(1)(d): criminal offence risk assessment based solely on profiling or personality traits.

Article 5(1)(e): untargeted facial-image scraping from the internet or CCTV footage.

Article 5(1)(f): workplace and education emotion recognition, except for medical or safety reasons.

Article 5(1)(g): **biometric** categorisation to infer sensitive characteristics.

Article 5(1)(h): real-time remote biometric identification in publicly accessible spaces for law enforcement, with narrow exceptions.



Table of contents

Search for content within t

Chapter I: General Provisions +

Chapter II: Prohibited AI Practices -

| Article 5: Prohibited AI Practices

Chapter III: High-Risk AI System +

Chapter IV: Transparency +

Copy URL

Part of [Chapter II: Prohibited AI Practices](#)

Article 5: Prohibited AI Practices

Date of entry into force: According to: Inherited from:

2 February 2025 **Article 113(a)** **Chapter II**

See here for a full implementation timeline.

SUMMARY +

1. The following AI practices shall be prohibited:

2. High-Risk AI Systems



Critical Infrastructure & Health

AI in water, gas, electricity, or medical devices that could put life and health at risk.



Education & Employment

Systems determining access to education or used for recruitment and workplace monitoring.



Law Enforcement & Justice

AI used in border control, administration of justice, and democratic processes.



Compliance Requirements

Subject to strict obligations: risk controls, documentation, monitoring, human oversight, and accountability mechanisms.

2. High-Risk AI Systems

Article 6: defines when an AI system is classified as high-risk.

Annex III: lists high-risk use areas, including biometrics, critical infrastructure, education, employment, essential services, law enforcement, migration, and justice.

Articles 9–15: set requirements for high-risk AI systems, including risk management, data governance, technical documentation, logging, transparency to deployers, human oversight, accuracy, robustness, and cybersecurity.

Part of [Chapter III: High-Risk AI System](#) → [Section 1: Classification of AI Systems as High-Risk](#)

Article 6: Classification Rules for High-Risk AI Systems

Date of entry into force:

2 August 2026 (2 August 2027 for Point 1)

See here for a [full implementation timeline](#).

According to:

Article 113

SUMMARY +

1. Irrespective of whether an AI system is placed on the market or put into service independently of the products referred to in points (a) and (b), that AI system shall be considered to be high-risk where both of the following conditions are fulfilled:

(a) the AI system is intended to be used as a safety component of a product, or the AI system is itself a product, covered by the Union harmonisation legislation listed in [Annex I](#);

(b) the product whose safety component pursuant to point (a) is the AI system, or the AI system itself as a product, is required to undergo a third-party conformity assessment, with a view to the placing on the market or the putting into service of that product pursuant to the Union harmonisation legislation listed in [Annex I](#).

2. High-Risk AI Systems

Annex III: lists high-risk use areas, including:

- Critical infrastructure
- Education and vocational training
- Employment and worker management
- Access to essential private and public services, including creditworthiness assessment
- Law enforcement
- Migration, asylum, and border control
- Administration of justice and democratic processes

High-risk AI systems pursuant to [Article 6\(2\)](#) are the AI systems listed in any of the following areas:

1. Biometrics, in so far as their use is permitted under relevant Union or national law:
 - (a) remote biometric identification systems. This shall not include AI systems intended to be used for biometric verification the sole purpose of which is to confirm that a specific natural person is the person he or she claims to be;
 - (b) AI systems intended to be used for biometric categorisation, according to sensitive or protected attributes or characteristics based on the inference of those attributes or characteristics;
 - (c) AI systems intended to be used for emotion recognition.

Related: [Recital 54](#) and [Recital 159](#)

2. Critical infrastructure: AI systems intended to be used as safety components in the management and operation of critical digital infrastructure, road traffic, or in the supply of water, gas, heating or electricity. Related: [Recital 55](#)

3. Education and vocational training:

- (a) AI systems intended to be used to determine access or admission or to assign natural persons to educational and vocational training institutions at all levels;
- (b) AI systems intended to be used to evaluate learning outcomes, including

2. High-Risk AI Systems

Articles 9–15: set requirements for high-risk AI systems, including: risk management, data governance, technical documentation, logging, transparency to deployers, human oversight, accuracy, robustness, and cybersecurity.

Section 2: Requirements for High-Risk AI Systems

Article 8: Compliance with the Requirements

Article 9: Risk Management System

Article 10: Data and Data Governance

Article 11: Technical Documentation

Article 12: Record-Keeping

Article 13: Transparency and Provision of Information to Deployers

Article 14: Human Oversight

Article 15: Accuracy, Robustness and Cybersecurity

1. A risk management system shall be established, implemented, documented and maintained in relation to high-risk AI systems.

2. The risk management system shall be understood as a continuous iterative process planned and run throughout the entire lifecycle of a high-risk AI system, requiring regular systematic review and updating. It shall comprise the following steps:

(a) the identification and analysis of the known and the reasonably foreseeable risks that the high-risk AI system can pose to health, safety or fundamental rights when the high-risk AI system is used in accordance with its intended purpose;

(b) the estimation and evaluation of the risks that may emerge when the high-risk AI system is used in accordance with its intended purpose, and under conditions of reasonably foreseeable misuse;

(c) the evaluation of other risks possibly arising, based on the analysis of data gathered from the post-market monitoring system referred to in [Article 72](#);

(d) the adoption of appropriate and targeted risk management measures designed to address the risks identified pursuant to point (a).

3. The risks referred to in this Article shall concern only those which may be reasonably mitigated or eliminated through the development or design of the high-risk AI system, or the provision of adequate technical information.

2. High-Risk AI Systems

Main takeaway

High-risk AI systems are regulated through **process obligations**, rather than outcome guarantees.

The law does not guarantee that a model will always be fair, accurate, or harmless. Instead, it requires responsible actors to demonstrate that they have systematically identified, documented, monitored, and mitigated risks.



Critical Infrastructure & Health

AI in water, gas, electricity, or medical devices that could put life and health at risk.



Education & Employment

Systems determining access to education or used for recruitment and workplace monitoring.



Law Enforcement & Justice

AI used in border control, administration of justice, and democratic processes.



Compliance Requirements

Subject to strict obligations: risk controls, documentation, monitoring, human oversight, and accountability mechanisms.

3. Transparency / Limited Risk AI

Key transparency obligations under Article 50

- Article 50(1): AI systems directly interacting with natural persons must inform people that they are interacting with an AI system.
- Article 50(2): Synthetic audio, image, video, or text: providers must ensure outputs are marked in a machine-readable way and detectable as artificially generated or manipulated.
- Article 50(3): Emotion recognition or biometric categorisation systems: affected people must be informed that such systems are being used.
- Article 50(4): Deepfakes: deployers must disclose that the image/audio/video content has been artificially generated or manipulated.
- AI-generated text on matters of public interest: deployers must disclose that the text was artificially generated or manipulated, unless it has gone through human review/editorial control.

Core

Limited-risk AI systems are generally allowed, but people must be informed about whether they are interacting with AI or seeing AI-generated/manipulated content.

3. Transparency / Limited Risk AI

Examples for discussion

- Chatbot: does the user know they are talking to AI?
- Deepfake: does the viewer know the video/audio is synthetic?
- AI-generated news-like text: is the public interest context clearly disclosed?
- Emotion recognition / biometric categorisation: do affected people know the system is being used?

Core

people must be **informed**
about whether they are
interacting with AI

4. Minimal / No Risk AI

Include: AI-enabled video games, spam filters.

- The EU AI Act does not impose new obligations on AI systems that are considered minimal or no risk.
- When the Act entered into force: minimal-risk systems face no obligations, though companies may voluntarily adopt codes of conduct.
- The Act is not a blanket regulation of all AI. Its central logic is proportionality: the higher the potential risk to health, safety, or fundamental rights, the stronger the legal obligations.

Discussion

**Where would you place each case in the EU AI Act risk pyramid?
Would your answer change depending on who deploys the system, what data it uses, or whether users are informed?**

Case 1: AI video interview system for hiring

A company uses an AI system to rank job candidates based on video interviews, speech patterns, language, and behavioral signals.

Case 2: AI spam filter in email

An email service uses AI to detect spam, phishing, and malware before they reach users' inboxes.

Case 3: Facial recognition database from scraped online images

A company scrapes billions of facial images from social media, news sites, and public websites to build a searchable facial recognition database used by law enforcement.

Case 4: AI-generated celebrity deepfake used in a scam

A fake AI-generated video of a celebrity appears on TikTok, falsely promoting a get-rich-quick scheme and encouraging users to sign up.

Discussion

Case 1: AI video interview / hiring ranking

High Risk.

Allowed, but subject to high-risk obligations: risk management, data governance, technical documentation, logging, human oversight, accuracy/robustness/cybersecurity; Annex III employment + Articles 9–15

Case 2: AI spam filter in email

Minimal / No Risk

No specific AI Act obligations; voluntary codes of conduct possible; ordinary consumer/data protection rules may still apply

Case 3: Scraped facial recognition database

Unacceptable Risk.

Prohibited if untargeted scraping of internet/CCTV facial images to create or expand facial recognition databases; Article 5(1)(e)

Case 4: AI-generated celebrity deepfake used in a scam

Limited / Transparency Risk: possibly illegal under other laws if fraud, harassment, or non-consensual sexual content
Disclosure / labeling obligations for deepfakes and AI-generated or manipulated content; Article 50(4); fraud or privacy harms may trigger other legal regimes

Discussion

Case 3: Facial recognition database from scraped online images

- A company scrapes billions of facial images from social media, news sites, and public websites to build a searchable facial recognition database used by law enforcement.
- -> Clearview AI has faced multiple European privacy enforcement actions for scraping facial images and building biometric databases. The Dutch Data Protection Authority fined Clearview €30.5 million in 2024 for illegal data collection for facial recognition

Home > Current >

Dutch DPA imposes a fine on Clearview because of illegal data collection for facial recognition

03 September 2024 Themes: [Biometrics](#) [Personal data on the Internet](#)

The Dutch Data Protection Authority (Dutch DPA) imposes a fine of 30.5 million euro and orders subject to a penalty for non-compliance up to more than 5 million euro on Clearview AI. Clearview is an American company that offers facial recognition services. Among other things, Clearview has built an illegal database with billions of photos of faces, including of Dutch people. The Dutch DPA warns that using the services of Clearview is also prohibited.

Clearview is a commercial business that offers facial recognition services to intelligence and investigative services. Customers of Clearview can provide



PUBLICATIONS

[Decision fine Clearview AI](#)

PDF, 1 MB

<https://www.autoriteitpersoonsgegevens.nl/en/current/dutch-dpa-imposes-a-fine-on-clearview-because-of-illegal-data-collection-for-facial-recognition>

Discussion

Case 4: AI-generated celebrity deepfake used in a scam

- A fake AI-generated video of a celebrity appears on TikTok, falsely promoting a get-rich-quick scheme and encouraging users to sign up
- → Recent deepfake scams have used fabricated celebrity videos, including Taylor Swift and other celebrities, to promote fake schemes and collect personal data..

New fake Taylor Swift TikTok scam reveals scary reality of AI deepfakes

News By Natalie Fear published 2 May 2026

No wonder the star is seeking legal action.



(Image credit: Disney Plus/ Taylor Swift)

<https://www.creativebloq.com/ai/new-fake-taylor-swift-tiktok-scam-reveals-scary-reality-of-ai-deepfakes>

Feedback from Public Consultation in 2020

Since first published, by 2023 the Act had more than 300 suggestions from different industries and sectors including automotive, pharmaceutical and technology.



Legal & Accountability Concerns

The Act was criticized for not accurately recognizing wrongs/harms or allocating responsibility. Experts argued it fell short of legally trustworthy AI standards.



SME Impact

Institutions warned of the risk of overburdening Small and Medium Enterprises (SMEs), citing concerns from the European Digital SME Alliance.



Strategic & Coordination Feedback (CLAIRE)

CLAIRE warned the regulation was too broad and vague to effectively boost "AI made in Europe" or strengthen EU competitiveness.

- Modest funding in relation to ambitious goals.
- Lack of coordination between instruments and mechanisms.
- Absence of large-scale signature initiatives to address fragmented ecosystem.

Economic & Regulatory Feedback (2020)



Projected Economic Impact

The Center for Data Innovation reported that the AI Act could cost the industry **€31 billion** over 5 years, potentially reducing investments by 20%.

**CEPS [12] suggests a significantly lower financial impact.*



Precautionary Principle Risks

Warning against early over-regulation: It is too soon to know how organizations and consumers will respond to AI developments.



OpenAI: Contextual Risk Assessment

OpenAI (creators of GPT-3) highlighted that a single AI application's risk level is not fixed, but varies significantly based on its specific usability and deployment context.

- Identifies "High-Risk" vs "Low-Risk" duality
- Emphasis on usability as the primary risk determinant

Feedback from Industry Leaders (2020)



Google

Identified the need to “hold providers and deployers to feasible standards” noting that some requirements are *“difficult or impossible to comply with in real-world situations,”* especially for **general-purpose AI systems**.



IBM

Argued that *“existing bodies are best placed”* to regulate high-risk AI, leveraging their specific sectoral expertise and established relationships with stakeholders.



Facebook (Meta)

Warned that the Commission’s approach could *“threaten to limit AI innovation across Europe”* and might inadvertently regulate technologies that aren't even AI systems.

Post GPT-3 effects on AI Act

Why GPAI Was Added

- The original AI Act was mainly designed around AI systems deployed in specific sectors.
- After the rise of generative AI, the final Act also had to regulate upstream models that can be [reused across many downstream contexts](#).
- Article 3(63): defines “general-purpose AI model” as a model that displays significant generality, can perform a wide range of distinct tasks, and can be integrated into a variety of downstream systems or applications.

What Counts as a GPAI Model?

Under the AI Act, a GPAI model is an AI model that:

- is trained at scale, often with large amounts of data
- displays significant generality
- can competently perform a wide range of distinct tasks
- can be integrated into many downstream systems or applications
- is not merely a research, development, or prototyping model before market placement
- **Examples:** Large language models; Multimodal foundation models; Image generation models

Who Should we regulate ?

Should We Regulate Models, Applications, or Both?

If one foundation model can power both a harmless chatbot and a high-risk hiring tool, where should legal responsibility begin?

Use-Case-Based Regulation

- Looks at where and how AI is deployed
- Fits the idea that “risk depends on context”
- Easier to connect to concrete harm
- May underestimate upstream model influence

Model-Level Regulation

- Looks at model capability, scale, and diffusion
- Fits the reality of foundation models reused everywhere
- Better captures upstream responsibility
- May overburden general-purpose technology or open research

Feedback from Public Consultation (2020)




The Need for Balanced Regulation

Regulating AI solely on risk metrics may neglect other medium-term risks. Over Regulating low-error applications can unreasonably affect European innovation rates.



Huawei: Multi-Actor Governance

Huawei identified that a *"one-size-fits-all"* approach fails to meet participant requirements. Advocacy for agent-based regulation accounting for different perspectives.



Implementation Timeline



Key Implementation Details

- The regulation spans 180 recitals and 113 Articles, totaling 144 pages of comprehensive rules.
- Phased rollout ensures stakeholders have time to adapt, starting with prohibited AI systems (6 months).
- GPAI rules and governance structures are prioritized to take effect within 12 months.

Final discussions

Is risk-based regulation enough? Can risk be classified in advance, or only in context?

Pros: avoids one-size-fits-all regulation; focuses attention on high-risk uses; balances innovation and fundamental rights.

Limits: categories may be rigid; the same model can be low-risk in one context and high-risk in another.

Can technical compliance solve social harm?

E.g., If a hiring AI system meets all compliance requirements, does that mean it is fair?

Tension: fairness is not only statistical; oversight can be symbolic; documentation may not give affected people real power.

Where is the risk in foundation models?

-> Is risk located in the model, the application, the deployer, or the social context?

Takeaway: The AI Act is also about assigning responsibility, not just regulating technology.

References

Council of the European Union. (2022, October 19). *Proposal for a regulation of the European Parliament and of the Council laying down harmonised rules on artificial intelligence (Artificial Intelligence Act) and amending certain Union legislative acts - Presidency compromise text.*
<https://artificialintelligenceact.eu/wp-content/uploads/2022/10/AIA-CZ-4th-Proposal-19-Oct-22.pdf>

European Commission. (2020) Have Your Say. Feedbacks on the proposal Artificial intelligence – ethical and legal requirements.
https://ec.europa.eu/info/law/better-regulation/have-your-say/initiatives/12527-Artificial-intelligence-ethical-and-legal-requirements/F2665626_en

European Commission. (2021). *Proposal for a regulation of the European Parliament and of the Council laying down harmonised rules on artificial intelligence (Artificial Intelligence Act) and amending certain Union legislative acts (COM/2021/206 final).* EUR-Lex.
<https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=celex%3A52021PC0206>

Future of Life Institute. (2022). *AI Act explorer.* EU Artificial Intelligence Act. Retrieved on the 30th April 2026, from
<https://artificialintelligenceact.eu/ai-act-explorer/>

Mueller, B. (2021). *How much will the Artificial Intelligence Act cost Europe?* Center for Data Innovation.
<https://www2.datainnovation.org/2021-aia-costs.pdf>

References

Renda, A. (2021, August 3). *Clarifying the costs for the EU's AI Act*. CEPS. <https://www.ceps.eu/clarifying-the-costs-for-the-eus-ai-act/>

White & Case LLP. (2024, July 15). *Long-awaited EU AI Act becomes law after publication in the EU's Official Journal*.

<https://www.whitecase.com/insight-alert/long-awaited-eu-ai-act-becomes-law-after-publication-eus-official-journal>

Veale, M., & Zuiderveen Borgesius, F. (2021). *Demystifying the draft EU Artificial Intelligence Act*. SSRN.

<https://doi.org/10.2139/ssrn.3899991>